

Python for Data Science

General Assembly, NY

Something to start with

With your best knowledge, write a python script that:

1. loads a file from a website

`(http://stat.columbia.edu/~rachel/datasets/nyt1.csv)`

2. from that file, counts the number of 1s and 0s under the 'Gender' column, and the number of 1s and 0s under the 'Signed_In' column

A Simple Solution

Read the script below.

```

## Download the file into ~/Downloads
## run: python nytimes_counter.py < ~/Downloads/nyt1.csv
# Import required libraries
import sys

# Start a counter and store the textfile in memory
gender = 0
signins = 0
lines = sys.stdin.readlines()
lines.pop(0)

# For each line, find the sum of index 2 in the list.
for line in lines:

    gender = gender + int(line.strip().split(',')[1])

for line in lines:
    signins = signins + int(line.strip().split(',')[4])

gender_0 = len(lines) - gender
signins_0 = len(lines) - signins

print "Gender 0: ", gender_0
print "Gender 1: ", gender_1
print "Signin 0: ", signins_0
print "Signin 1: ", signins_1

```

Optimizations

Download in terminal, or python

On unix machines, we can use either curl or wget to download the file:

```

curl http://stat.columbia.edu/~rachel/datasets/nyt1.csv > nyt1.csv
wget http://stat.columbia.edu/~rachel/datasets/nyt1.csv > nyt1.csv

```

on PC, you can use powershell:

```
Invoke-WebRequest http://stat.columbia.edu/~rachel/datasets/nyt1.csv -OutFile  
nyt1.csv
```

since we're going all the data manipulation in python, we could also use a library to store the data in memory. Below are two common approaches to this.

```
import csv  
import requests  
import urllib2  
import StringIO  
  
url = 'http://stat.columbia.edu/~rachel/datasets/nyt1.csv'  
  
## urllib2 version  
response = urllib2.urlopen(url)  
nyt = csv.reader(response)  
  
## requests version  
r = requests.get(url)  
data = r.text  
nyt = csv.reader(data.splitlines(), delimiter='\t')
```

Loop once if you only have to loop once

One way to improve the script above is to loop through the iterator `nyt` only once:

```
for line in nyt:  
    gender = gender + line[1]  
    signins = signins + line[4]
```

New Script

```
In [2]: import csv
import requests
import urllib2
import StringIO

url = 'http://stat.columbia.edu/~rachel/datasets/nyt1.csv'

response = urllib2.urlopen(url)
nyt = csv.reader(response)

counts, gender, signins = 0, 0, 0

# ignores the header row
next(nyt, None)
for line in nyt:
    counts += 1
    gender += int(line[1])
    signins += int(line[4])

gender_0 = counts - gender
signins_0 = counts - signins

print "Gender 0:", gender_0
print "Gender 1:", gender
print "Signin 0:", signins_0
print "Signin 1:", signins

Gender 0: 290176
Gender 1: 168265
Signin 0: 137106
Signin 1: 321335
```

Learning Python in an Hour

adapted from Alysaa Frazee, [introducing R to a non-programmer in one hour](http://alyssafrazee.com/introducing-R.html)
(<http://alyssafrazee.com/introducing-R.html>)

```

In [19]: x = 7
         print x + 5
         # These are comments! And they are super helpful!
         # use help() to get help about what something is doing
         help(x)

12
Help on int object:

class int(object)
|   int(x=0) -> int or long
|   int(x, base=10) -> int or long
|
|   Convert a number or string to an integer, or return 0
if no arguments
|   are given.  If x is floating point, the conversion tr
uncates towards zero.
|   If x is outside the integer range, the function retur
ns a long instead.
|
|   If x is not a number or if base is given, then x must
be a string or
|   Unicode object representing an integer literal in the
given base.  The
|   literal can be preceded by '+' or '-' and be surround
ed by whitespace.
|   The base defaults to 10.  Valid bases are 0 and 2-36.
Base 0 means to
|   interpret the base from the string as an integer lite
ral.
|   >>> int('0b100', base=0)
|   4
|
|   Methods defined here:
|
|   __abs__(...)
|       x.__abs__() <==> abs(x)
|
|   __add__(...)
|       x.__add__(y) <==> x+y
|
|   __and__(...)
|       x.__and__(y) <==> x&y
|
|   cmp    (...)

```

```

|   __cmp__(y) <==> cmp(x, y)
|
|   __coerce__(...)
|
|   x.__coerce__(y) <==> coerce(x, y)
|
|   __div__(...)
|   x.__div__(y) <==> x/y
|
|   __divmod__(...)
|   x.__divmod__(y) <==> divmod(x, y)
|
|   __float__(...)
|   x.__float__() <==> float(x)
|
|   __floordiv__(...)
|   x.__floordiv__(y) <==> x//y
|
|   __format__(...)
|
|   __getattr__(...)
|   x.__getattr__('name') <==> x.name
|
|   __getnewargs__(...)
|
|   __hash__(...)
|   x.__hash__() <==> hash(x)
|
|   __hex__(...)
|   x.__hex__() <==> hex(x)
|
|   __index__(...)
|   x[y:z] <==> x[y.__index__() : z.__index__()]
|
|   __int__(...)
|   x.__int__() <==> int(x)
|
|   __invert__(...)
|   x.__invert__() <==> ~x
|
|   __long__(...)
|   x.__long__() <==> long(x)
|
|   __lshift__(...)
|   x.__lshift__(y) <==> x<<y

```

```

__mod__(...)
    x.__mod__(y) <==> x%y

__mul__(...)
    x.__mul__(y) <==> x*y

__neg__(...)
    x.__neg__() <==> -x

__nonzero__(...)
    x.__nonzero__() <==> x != 0

__oct__(...)
    x.__oct__() <==> oct(x)

__or__(...)
    x.__or__(y) <==> x|y

__pos__(...)
    x.__pos__() <==> +x

__pow__(...)
    x.__pow__(y[, z]) <==> pow(x, y[, z])

__radd__(...)
    x.__radd__(y) <==> y+x

__rand__(...)
    x.__rand__(y) <==> y&x

__rdiv__(...)
    x.__rdiv__(y) <==> y/x

__rdivmod__(...)
    x.__rdivmod__(y) <==> divmod(y, x)

__repr__(...)
    x.__repr__() <==> repr(x)

__rfloordiv__(...)
    x.__rfloordiv__(y) <==> y//x

__rlshift__(...)
    x.__rlshift__(y) <==> y<<x

```

```

__rmod__(...)
    x.__rmod__(y) <==> y%x

__rmul__(...)
    x.__rmul__(y) <==> y*x

__ror__(...)
    x.__ror__(y) <==> y|x

__rpow__(...)
    y.__rpow__(x[, z]) <==> pow(x, y[, z])

__rrshift__(...)
    x.__rrshift__(y) <==> y>>x

__rshift__(...)
    x.__rshift__(y) <==> x>>y

__rsub__(...)
    x.__rsub__(y) <==> y-x

__rtruediv__(...)
    x.__rtruediv__(y) <==> y/x

__rxor__(...)
    x.__rxor__(y) <==> y^x

__str__(...)
    x.__str__() <==> str(x)

__sub__(...)
    x.__sub__(y) <==> x-y

__truediv__(...)
    x.__truediv__(y) <==> x/y

__trunc__(...)
    Truncating an Integral returns itself.

__xor__(...)
    x.__xor__(y) <==> x^y

bit_length(...)
    int.bit_length() -> int

```



```

|         Number of bits necessary to represent self in binary.
|
|         >>> bin(37)
|
|         '0b100101'
|         >>> (37).bit_length()
|         6
|
|         conjugate(...)
|             Returns self, the complex conjugate of any int.
|
|         -----
|
|         Data descriptors defined here:
|
|         denominator
|             the denominator of a rational number in lowest terms
|
|         imag
|             the imaginary part of a complex number
|
|         numerator
|             the numerator of a rational number in lowest terms
|
|         real
|             the real part of a complex number
|
|         -----
|
|         Data and other attributes defined here:
|
|         __new__ = <built-in method __new__ of type object>
|             T.__new__(S, ...) -> a new object with type S, a
|             subtype of T

```

Basic Data Types

Dictionaries, functions

```
In [45]: some_dictionary = dict()
some_dictionary['fruits'] = ['apples', 'oranges', 'bananas']
some_dictionary['veggies'] = ['beans', 'carrots', 'kale', 'beats']
print some_dictionary['fruits']
print some_dictionary.keys()
print some_dictionary.values()
print some_dictionary.items()
print

def print_items(dictionary):
    for k,v in dictionary.iteritems():
        print '%s: %s' % (k, v,)

def longest_list(dictionary):
    return max(enumerate(dictionary.values()), key = lambda tup: len(tup[1]))

print_items(some_dictionary)

print longest_list(some_dictionary)

['apples', 'oranges', 'bananas']
['veggies', 'fruits']
[['beans', 'carrots', 'kale', 'beats'], ['apples', 'oranges', 'bananas']]
[('veggies', ['beans', 'carrots', 'kale', 'beats']), ('fruits', ['apples', 'oranges', 'bananas'])]

veggies: ['beans', 'carrots', 'kale', 'beats']
fruits: ['apples', 'oranges', 'bananas']
(0, ['beans', 'carrots', 'kale', 'beats'])
```

Classes

Python is considered Object Oriented Programming (OOP), and because of that, much of the functionality with libraries used in Python and in the Data Science course will be Classes. Consider the following changes to the script above:

```
In [185]: import csv
import requests
import urllib2
import StringIO

class NYTimesCounter():

    def __init__(self):
        # initializes the object
        self.counts = 0
        self.gender = 0
        self.signin = 0

    def reset(self):
        # Resets the counters
        self.counts = 0
        self.gender = 0
        self.signin = 0

    def _get_url(self, url):
        # private function that retrieves the file,
        # and if the url is local, retrieves locally i
        nstead.
        if url[0:4] == 'http':
            self.response = urllib2.urlopen(url)
        else:
            self.response = open(url)

    def run_file(self, url):
        # calls get_url() and counts the data
        self.reset()
        self._get_url(url)
        nyt = csv.reader(self.response)
        next(nyt, None)
        for line in nyt:
            self.counts += 1
            self.gender += int(line[1])
            self.signin += int(line[4])
```

```

    def print_params(self):
        print "Gender 0:", self.gender
        print "Gender 1:", self.counts - self.gender
        print "Signin 0:", self.signin
        print "Signin 1:", self.counts - self.gender

nyt = NYTimesCounter()
print type(nyt)
#nyt.run_file('http://stat.columbia.edu/~rachel/dataset/nyt1.csv')
nyt.run_file('/Users/edjoy/Downloads/nyt1.csv')
nyt.print_params()

<type 'instance'>
Gender 0: 168265
Gender 1: 290176
Signin 0: 321335
Signin 1: 290176

```

Takeaways

1. How is the above script different from what we ran previously?
2. What advantages and disadvantages do you see in OOP vs functional programming (which is similar to the earlier function)

Libraries for Data Science

Data Scientists use a wide variety of libraries in Python that make working with data significantly easier. Those libraries primarily consist of:

1. `numpy`
2. `scipy`
3. `pandas`
4. `matplotlib`
5. `statsmodels`
6. `scikit-learn`
7. `nltk`

Though there are countless others available.

For today, we'll primarily focus ourselves around `pandas`, `matplotlib`, and `sklearn`.

pandas

`pandas` is a library built on top of `numpy`, which allows us to use excel-like matrices in the python programming space. These special matrices are called DataFrames.

Load up the `nyt1.csv` (even from the website directly) as a `pandas DataFrame` like so:

```
In [186]: import pandas as pd
```

```
#nyt = pd.read_csv('http://stat.columbia.edu/~rachel/datasets/nyt1.csv')  
nyt = pd.read_csv('/Users/edjoy/Downloads/nyt1.csv')
```

```
In [187]: print type(nyt)
          print nyt.dtypes

          print nyt.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Age          int64
Gender        int64
Impressions   int64
Clicks        int64
Signed_In     int64
dtype: object
```

	Age	Gender	Impressions	Clicks \
count	458441.000000	458441.000000	458441.000000	458441.000000
mean	29.482551	0.367037	5.007316	0.092594
std	23.607034	0.481997	2.239349	0.309973
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	3.000000	0.000000
50%	31.000000	0.000000	5.000000	0.000000
75%	48.000000	1.000000	6.000000	0.000000
max	108.000000	1.000000	20.000000	4.000000

	Signed_In
count	458441.000000
mean	0.700930
std	0.457851
min	0.000000
25%	0.000000
50%	1.000000
75%	1.000000
max	1.000000

The describe function for a data frame creates a high level view of what your data looks like. With quantile distributions plus a mean (average), We can measure the same information we had before ($.701 * 458441 = \sim 321360$ Gender 0), though we can use other built in functions for this as well.

Like dictionaries, you call on columns using their keys, and like lists, you can subset on indices.

```
In [3]: print nyt['Gender'].sum()
        print nyt['Signed_In'].sum()

        print nyt[1:3]
        print nyt.head()
        print nyt.head(10)

        print nyt.tail()
```


168265

321335

	Age	Gender	Impressions	Clicks	Signed_In
1	73	1	3	0	1
2	30	0	3	0	1

	Age	Gender	Impressions	Clicks	Signed_In
0	36	0	3	0	1
1	73	1	3	0	1
2	30	0	3	0	1
3	49	1	3	0	1
4	47	1	11	0	1

	Age	Gender	Impressions	Clicks	Signed_In
0	36	0	3	0	1
1	73	1	3	0	1
2	30	0	3	0	1
3	49	1	3	0	1
4	47	1	11	0	1
5	47	0	11	1	1
6	0	0	7	1	0
7	46	0	5	0	1
8	16	0	3	0	1
9	52	0	4	0	1

	Age	Gender	Impressions	Clicks	Signed_In
458436	0	0	2	0	0
458437	0	0	4	0	0
458438	72	1	5	0	1
458439	0	0	5	0	0
458440	0	0	3	0	0

DataFrames allow you to subset as well. What does the data look like if you subset this click and impression data by Signed_In?

```
In [4]: print nyt[nyt['Signed_In'] == 0].describe()
```

groupby is also an effective way to create pivot tables--but here we just use it as a simpler way to see both data segmentations

```
print nyt.groupby('Signed_In').describe()
```

	Age	Gender	Impressions	Clicks	Signed_In
count	137106	137106	137106.000000	137106.000000	13

7106				
mean	0	0	4.999657	0.14208
0				
std	0	0	2.240662	0.38551
0				
min	0	0	0.000000	0.00000
0				
25%	0	0	3.000000	0.00000
0				
50%	0	0	5.000000	0.00000
0				
75%	0	0	6.000000	0.00000
0				
max	0	0	18.000000	4.00000
0				

		Age	Clicks	Gen
der	Impressions			
Signed_In				
0	count	137106.000000	137106.000000	137106.000
000	137106.000000			
	mean	0.000000	0.142080	0.000
000	4.999657			
	std	0.000000	0.385510	0.000
000	2.240662			
	min	0.000000	0.000000	0.000
000	0.000000			
	25%	0.000000	0.000000	0.000
000	3.000000			
	50%	0.000000	0.000000	0.000
000	5.000000			
	75%	0.000000	0.000000	0.000
000	6.000000			
	max	0.000000	4.000000	0.000
000	18.000000			
1	count	321335.000000	321335.000000	321335.000
000	321335.000000			
	mean	42.062054	0.071480	0.523
644	5.010584			
	std	16.308117	0.268659	0.499
441	2.238784			
	min	7.000000	0.000000	0.000
000	0.000000			
	25%	29.000000	0.000000	0.000
000	3.000000			

	50%	41.000000	0.000000	1.000
000	5.000000			
	75%	53.000000	0.000000	1.000
000	6.000000			
	max	108.000000	3.000000	1.000
000	20.000000			

NEXT STEPS

What's the objective? We want to see if there is a correlation with age, gender, and Signed_Out with click_thru_rate, measured as clicks / Impressions

other ways to fiddle around with pandas (slicing) create a function that defines click through rate for each row

introduction to matplotlib

plot each and everything

```
In [5]: # Observe what occurs if we divide ints, compared to dividing floats:
        print 1 / 2
        print 1.0 / 2.0

0
0.5
```

```
In [191]: # in order to create click thru, we need clicks and im  
          pressions to be floats, otherwise they do not divide a  
          s a human would expect!
```

```
nyt['Clicks'] = nyt['Clicks'].astype('float')  
nyt['Impressions'] = nyt['Impressions'].astype('float'  
)
```

```
# You could also change them at the same time  
nyt[['Clicks', 'Impressions']] = nyt[['Clicks', 'Impre  
ssions']].astype('float')
```

```
# Or pass the columns in as a list variable
```

```
columns_to_float=['Clicks', 'Impressions']  
nyt[columns_to_float] = nyt[columns_to_float].astype(''  
float')
```

```
nyt['Click_Thru'] = nyt['Clicks'] / nyt['Impressions']
```

```
In [192]: print nyt.describe()  
          print nyt.groupby('Signed_In').describe()
```

	Age	Gender	Impressions	
Clicks \				
count	458441.000000	458441.000000	458441.000000	458441.000000
mean	29.482551	0.367037	5.007316	0.092594
std	23.607034	0.481997	2.239349	0.309973
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	3.000000	0.000000
50%	31.000000	0.000000	5.000000	0.000000
75%	48.000000	1.000000	6.000000	0.000000
max	108.000000	1.000000	20.000000	4.000000

	Signed_In	Click_Thru
count	458441.000000	455375.000000

mean	0.700930	0.018471
std	0.457851	0.069034
min	0.000000	0.000000
25%	0.000000	0.000000
50%	1.000000	0.000000
75%	1.000000	0.000000
max	1.000000	1.000000

	Age	Click_Thru	Cli
cks			
Signed_In	Gender \		

0	count	137106.000000	136177.000000	137106.000
000	137106.000000			
	mean	0.000000	0.028355	0.142
080	0.000000			
	std	0.000000	0.085324	0.385
510	0.000000			
	min	0.000000	0.000000	0.000
000	0.000000			
	25%	0.000000	0.000000	0.000
000	0.000000			
	50%	0.000000	0.000000	0.000
000	0.000000			
	75%	0.000000	0.000000	0.000
000	0.000000			
	max	0.000000	1.000000	4.000
000	0.000000			
1	count	321335.000000	319198.000000	321335.000
000	321335.000000			
	mean	42.062054	0.014254	0.071
480	0.523644			
	std	16.308117	0.060280	0.268
659	0.499441			
	min	7.000000	0.000000	0.000
000	0.000000			
	25%	29.000000	0.000000	0.000
000	0.000000			
	50%	41.000000	0.000000	0.000
000	1.000000			
	75%	53.000000	0.000000	0.000
000	1.000000			
	max	108.000000	1.000000	3.000
000	1.000000			

Impressions

Signed_In	
0	count 137106.000000
	mean 4.999657
	std 2.240662
	min 0.000000
	25% 3.000000
	50% 5.000000
	75% 6.000000
	max 18.000000
1	count 321335.000000
	mean 5.010584
	std 2.238784
	min 0.000000
	25% 3.000000
	50% 5.000000
	75% 6.000000
	max 20.000000

In [193]: # Keep in mind you can use lists in a group by as well
print nyt.groupby(['Signed_In', 'Gender']).describe()

			Age	Click_Thru	
Clicks \					
Signed_In Gender					
0	0	count	137106.000000	136177.000000	137
106.000000		mean	0.000000	0.028355	
0.142080		std	0.000000	0.085324	
0.385510		min	0.000000	0.000000	
0.000000		25%	0.000000	0.000000	
0.000000		50%	0.000000	0.000000	
0.000000		75%	0.000000	0.000000	
0.000000		max	0.000000	1.000000	
4.000000					
1	0	count	153070.000000	152052.000000	153
070.000000		mean	43.423336	0.014622	

0.073117				
	std	16.763906	0.060956	
0.271194				
	min	7.000000	0.000000	
0.000000				
	25%	30.000000	0.000000	
0.000000				
	50%	42.000000	0.000000	
0.000000				
	75%	55.000000	0.000000	
0.000000				
	max	108.000000	1.000000	
3.000000				
1	count	168265.000000	167146.000000	168
265.000000				
	mean	40.823701	0.013919	
0.069991				
	std	15.780505	0.059656	
0.266324				
	min	7.000000	0.000000	
0.000000				
	25%	28.000000	0.000000	
0.000000				
	50%	40.000000	0.000000	
0.000000				
	75%	52.000000	0.000000	
0.000000				
	max	107.000000	1.000000	
3.000000				

		Impressions		
Signed_In	Gender			
0	0	count	137106.000000	
		mean	4.999657	
		std	2.240662	
		min	0.000000	
		25%	3.000000	
		50%	5.000000	
		75%	6.000000	
1	0	max	18.000000	
		count	153070.000000	
		mean	5.012733	
		std	2.238426	
		min	0.000000	
		25%	3.000000	

	50%	5.000000
	75%	6.000000
	max	17.000000
1	count	168265.000000
	mean	5.008629
	std	2.239114
	min	0.000000
	25%	3.000000
	50%	5.000000
	75%	6.000000
	max	20.000000

matplotlib

matplotlib's core functionality serves as a plotting tool within python. While calling `.describe()` on DataFrames is useful to get a rough idea of what your data looks like, plots allow you to visualize what your data really looks like.

Consider the following data set and code:


```
In [194]: anscombe = pd.DataFrame({
    'x' : [10, 8, 13, 9, 11, 14, 6, 4, 12, 7, 5],
    'y1' : [8.04, 6.95, 7.58, 8.81, 8.33, 9.96, 7.24,
4.26, 10.84, 4.82, 5.68],
    'y2' : [9.14, 8.14, 8.74, 8.77, 9.26, 8.10, 6.13,
3.10, 9.13, 7.26, 4.74],
    'y3' : [7.46, 6.77, 12.74, 7.11, 7.81, 8.84, 6.08,
5.39, 8.15, 6.42, 5.73],
    'x4' : [8,8,8,8,8,8,8,19,8,8,8],
    'y4' : [6.58,5.76,7.71,8.84,8.47,7.04,5.25,12.50,5
.56,7.91,6.89],
})
```

```
anscombe.describe()
```

Out[194]:

	x	x4	y1	y2	y3	y4
count	11.000000	11.000000	11.000000	11.000000	11.000000	11.0000
mean	9.000000	9.000000	7.500909	7.500909	7.500000	7.50090
std	3.316625	3.316625	2.031568	2.031657	2.030424	2.03057
min	4.000000	8.000000	4.260000	3.100000	5.390000	5.25000
25%	6.500000	8.000000	6.315000	6.695000	6.250000	6.17000
50%	9.000000	8.000000	7.580000	8.140000	7.110000	7.04000
75%	11.500000	8.000000	8.570000	8.950000	7.980000	8.19000
max	14.000000	19.000000	10.840000	9.260000	12.740000	12.5000

Visually from creating the data frame you can tell the data looks different, yet in the `.describe()` call, the data shares very similar features. The two primary plotting tools we use from matplotlib are histograms and scatterplots, which help us understand the shape of data.

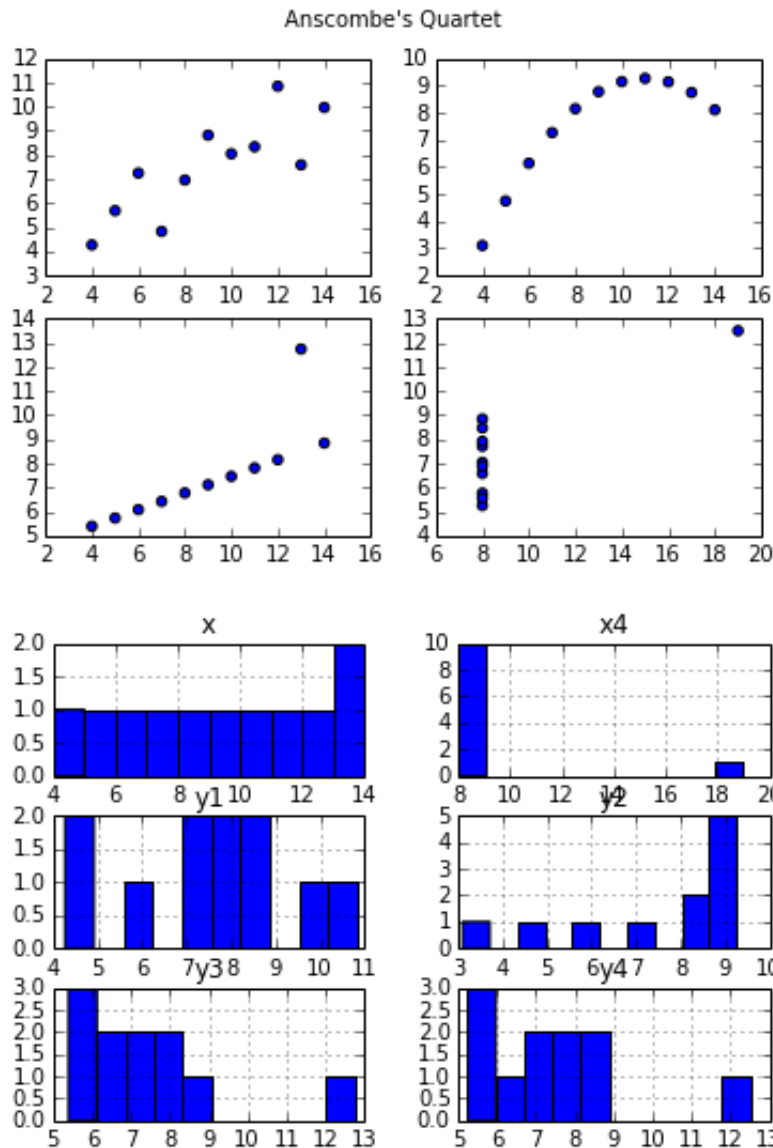
```
In [195]: from matplotlib import pylab as plt
```

```
fig = plt.figure()
ax1 = fig.add_subplot(2,2,1) # Two rows, Two columns,
First Plot
ax1.scatter(anscombe.x, anscombe.y1)
ax2 = fig.add_subplot(2,2,2) # Second Plot (Plots go L
-> R, Up -> Down)
ax2.scatter(anscombe.x, anscombe.y2)
ax3 = fig.add_subplot(2,2,3) # Third
ax3.scatter(anscombe.x, anscombe.y3)
ax4 = fig.add_subplot(2,2,4) # Fourth
ax4.scatter(anscombe.x4, anscombe.y4)
fig.suptitle("Anscombe's Quartet")
print fig

# Pandas also has matplotlib built in, which is incred
ibly useful for creating fast histograms.
print anscombe.hist()
```

Figure (480x320)

```
[(<matplotlib.axes.AxesSubplot object at 0x10b9b1950>  
  <matplotlib.axes.AxesSubplot object at 0x1108df110>]  
[(<matplotlib.axes.AxesSubplot object at 0x10dafce90>  
  <matplotlib.axes.AxesSubplot object at 0x10bd07950>]  
[(<matplotlib.axes.AxesSubplot object at 0x10bb874d0>  
  <matplotlib.axes.AxesSubplot object at 0x10baa8550>)]
```

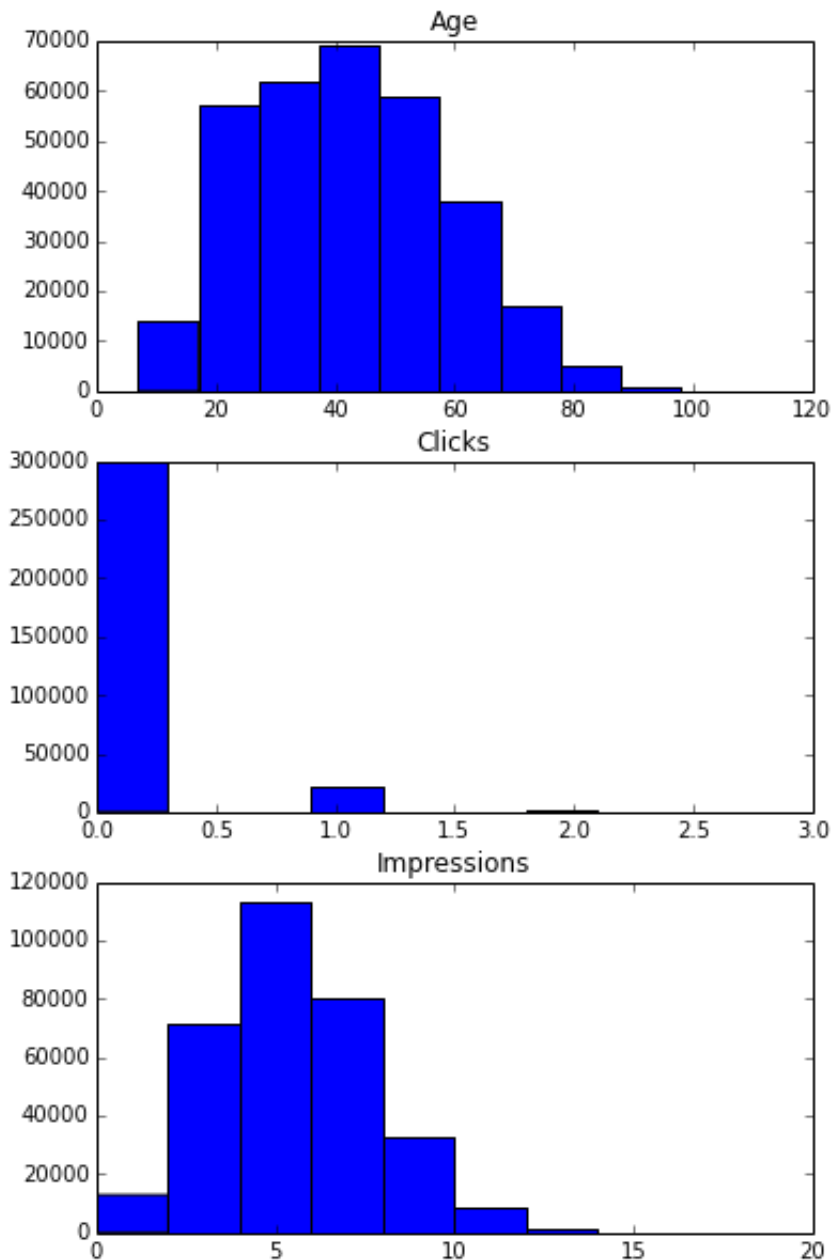


Use these new tools in order to visualize some of this New York Times ad performance data.

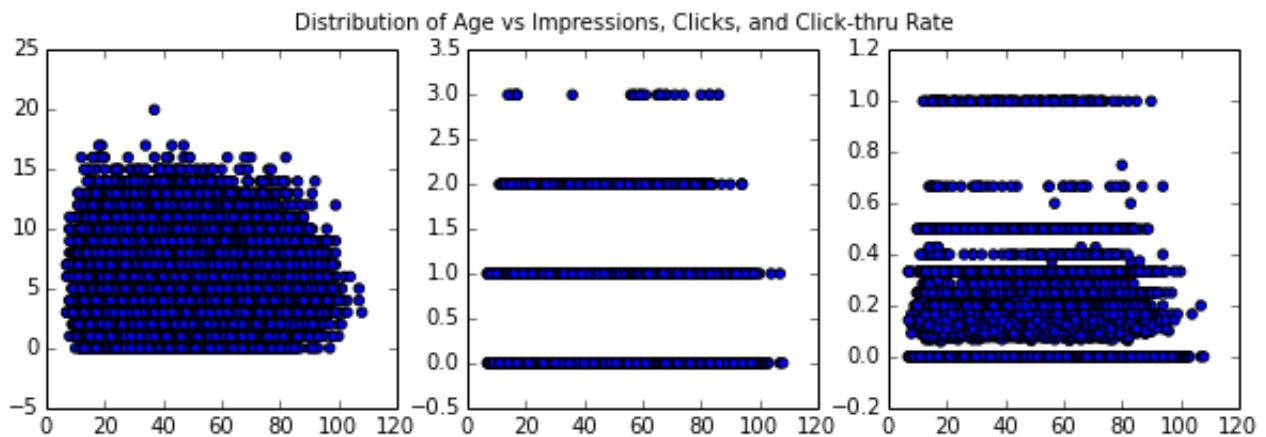
```
In [196]: nyt_signed_in = nyt[nyt['Signed_In'] ==1]
```

```
fig = plt.figure(figsize=(6, 10))  
ax1 = fig.add_subplot(3, 1, 1)
```

```
ax1.set_title('Age')  
ax1.hist(nyt_signed_in['Age'])  
ax2 = fig.add_subplot(3, 1, 2)  
ax2.set_title('Clicks')  
ax2.hist(nyt_signed_in['Clicks'])  
ax3 = fig.add_subplot(3, 1, 3)  
ax3.set_title('Impressions')  
ax3.hist(nyt_signed_in['Impressions'])  
fig.show()
```




```
In [197]: fig = plt.figure(figsize=(10, 3))
fig.suptitle('Distribution of Age vs Impressions, Clicks, and Click-thru Rate')
ax1 = fig.add_subplot(131)
ax1.scatter(nyt_signed_in['Age'], nyt_signed_in['Impressions'])
ax2 = fig.add_subplot(132)
ax2.scatter(nyt_signed_in['Age'], nyt_signed_in['Clicks'])
ax3 = fig.add_subplot(133)
ax3.scatter(nyt_signed_in['Age'], nyt_signed_in['Click_Thru_Rate'])
fig.show()
```



Even though we see some clear normal-like distributions [Why do we care about normal distributions?](http://www.quora.com/Normal-Distribution-statistics/Why-do-we-use-the-normal-distribution) (<http://www.quora.com/Normal-Distribution-statistics/Why-do-we-use-the-normal-distribution>) with histograms, Comparing Age to a few variables does not show any clear relationships.

```
In [198]: import numpy as np
# Just in case, let's compare just one gender:
nyt_gender0 = nyt_signed_in[nyt_signed_in['Gender'] ==
0]
nyt_gender0 = nyt_gender0[np.isfinite(nyt_gender0['Click_Thru'])]
fig = plt.figure(figsize=(18, 3))
fig.suptitle('Distribution of Age vs Impressions, Clicks, and Click-thru Rate')
ax1 = fig.add_subplot(141)
ax1.scatter(nyt_gender0['Age'], nyt_gender0['Impressions'])

ax2 = fig.add_subplot(142)
ax2.scatter(nyt_gender0['Age'], nyt_gender0['Clicks'])
ax3 = fig.add_subplot(143)
ax3.scatter(nyt_gender0['Age'], nyt_gender0['Click_Thru'])
ax4 = fig.add_subplot(144)
ax4.hist(nyt_gender0['Click_Thru'])
fig.show()
```



scikit-learn

Scikit-learn (often 'sklearn') is one of several core machine learning packages available in python.

scikit-learn is designed to be modular. Many parts of the libraries are super classes of base packages, which means that many of them share the same functionality. For example, consider making the following classes, Array and Matrix, where Matrix is actually a special kind of Array:

```
In [199]: class MyArray():
```

```
    def __init__(self, x, y, dim, value):
        self.x = x
        self.y = y
        self.dim = dim
        self.value = value

    def multi(self):
        values = []
        for i in range(self.x):
            values.append([self.value for k in xrange
(self.y)] for j in xrange(self.dim)])
        return values

    def show(self):
        print "Showing", self.__class__.__name__
        for i in self.multi():
            print i
```

```
my_array = MyArray(5, 8, 5, 0)
my_array.show()
```

```
print
```

```
class MyMatrix(MyArray):
    def __init__(self, x, y, value):
        self.x = x
        self.y = y
        self.dim = 1
        self.value = value
```

```
my_matrix = MyMatrix(4, 8, 0)
my_matrix.show()
```

Showing MyArray

```
[[0, 0, 0, 0, 0, 0, 0, 0], [0, 0, 0, 0, 0, 0, 0, 0], [0,
0, 0, 0, 0, 0, 0, 0], [0, 0, 0, 0, 0, 0, 0, 0], [0, 0, 0,
0, 0, 0, 0, 0]]
[[0, 0, 0, 0, 0, 0, 0, 0], [0, 0, 0, 0, 0, 0, 0, 0], [0,
0, 0, 0, 0, 0, 0, 0], [0, 0, 0, 0, 0, 0, 0, 0], [0, 0, 0,
0, 0, 0, 0, 0]]
[[0, 0, 0, 0, 0, 0, 0, 0], [0, 0, 0, 0, 0, 0, 0, 0], [0,
0, 0, 0, 0, 0, 0, 0], [0, 0, 0, 0, 0, 0, 0, 0], [0, 0, 0,
0, 0, 0, 0, 0]]
[[0, 0, 0, 0, 0, 0, 0, 0], [0, 0, 0, 0, 0, 0, 0, 0], [0,
0, 0, 0, 0, 0, 0, 0], [0, 0, 0, 0, 0, 0, 0, 0], [0, 0, 0,
0, 0, 0, 0, 0]]
[[0, 0, 0, 0, 0, 0, 0, 0], [0, 0, 0, 0, 0, 0, 0, 0], [0,
0, 0, 0, 0, 0, 0, 0], [0, 0, 0, 0, 0, 0, 0, 0], [0, 0, 0,
0, 0, 0, 0, 0]]
```

Showing MyMatrix

```
[[0, 0, 0, 0, 0, 0, 0, 0]]
[[0, 0, 0, 0, 0, 0, 0, 0]]
[[0, 0, 0, 0, 0, 0, 0, 0]]
[[0, 0, 0, 0, 0, 0, 0, 0]]
```

Notice how the MyMatrix class borrows functions from the MyArray class? The only difference is that Matrices are special kinds of Arrays, in that they are only 2 dimensional.

(note, a better way to do this would be to use the `numpy.array()`)

`scikit-learn` is in fact, very similar, as most functionality is built on a primary `.fit()` function for each learning algorithm (even their dummy regression, which is described [here](http://scikit-learn.org/stable/modules/generated/sklearn.dummy.DummyRegressor.html) (<http://scikit-learn.org/stable/modules/generated/sklearn.dummy.DummyRegressor.html>), follows the same base Class!)


```
In [200]: from sklearn.dummy import DummyRegressor
from sklearn.linear_model import LinearRegression

nyt = nyt[np.isfinite(nyt['Click_Thru'])]
dummy_fit = DummyRegressor().fit(nyt[['Signed_In', 'Gender', 'Age']], nyt['Click_Thru'])
lm_fit = LinearRegression().fit(nyt[['Signed_In', 'Gender', 'Age']], nyt['Click_Thru'])
anscombe_fit1 = LinearRegression().fit(anscombe[['x']], anscombe['y1'])
anscombe_fit2 = LinearRegression().fit(anscombe[['x']], anscombe['y2'])
anscombe_fit3 = LinearRegression().fit(anscombe[['x']], anscombe['y3'])
anscombe_fit4 = LinearRegression().fit(anscombe[['x4']], anscombe['y4'])
```

```
In [201]: print "Another point of weakness: Anscombe's Quartet has the same linear model for each group"
print anscombe_fit1.coef_, anscombe_fit1.intercept_
print anscombe_fit2.coef_, anscombe_fit2.intercept_
print anscombe_fit3.coef_, anscombe_fit3.intercept_
print anscombe_fit4.coef_, anscombe_fit4.intercept_

print
print dummy_fit.score(nyt[['Signed_In', 'Gender', 'Age']], nyt['Click_Thru'])
print lm_fit.score(nyt[['Signed_In', 'Gender', 'Age']], nyt['Click_Thru'])
```

Another point of weakness: Anscombe's Quartet has the same linear model for each group

```
[ 0.50009091] 3.00009090909
[ 0.5] 3.00090909091
[ 0.49972727] 3.00245454545
[ 0.49990909] 3.00172727273
```

0.0

0.0103350341913

Above, the `.score()` function, another common function across many `scikit-learn` Classes, returns back the r-squared value for this linear model. Unfortunately, it does not seem like you can fit the `nytimes` data linearly, as the score is very close to 0.

Datasets available within `scikit-learn`

There are also several data sets available within `scikit-learn`, however they primarily rely on the `numpy` objects instead. Become familiar with Fisher's Iris Data set, as it's a fairly common and unique data set, and will likely be referenced in the GA's primary data science class, or other classes you make take in the future.

```
In [153]: from sklearn import datasets
```

```
iris = datasets.load_iris()
```

```
print iris.DESCR
```

```
Iris Plants Database
```

```
Notes
```

```
-----
```

```
Data Set Characteristics:
```

```
    :Number of Instances: 150 (50 in each of three classe  
s)
```

```
    :Number of Attributes: 4 numeric, predictive attribut  
es and the class
```

```
    :Attribute Information:
```

- sepal length in cm
- sepal width in cm
- petal length in cm
- petal width in cm
- class:

- Iris-Setosa
- Iris-Versicolour
- Iris-Virginica

```
    :Summary Statistics:
```

```
=====
```

=====

	Min	Max	Mean	SD	Class Correlation
--	-----	-----	------	----	-------------------

=====

=====

sepal length:	4.3	7.9	5.84	0.83	0.7826	
sepal width:	2.0	4.4	3.05	0.43	-0.4194	
petal length:	1.0	6.9	3.76	1.76	0.9490	(high!)
petal width:	0.1	2.5	1.20	0.76	0.9565	(high!)

=====

=====

:Missing Attribute Values: None
:Class Distribution: 33.3% for each of 3 classes.
:Creator: R.A. Fisher
:Donor: Michael Marshall (MARSHALL%PLU@io.arc.nasa.gov)
:Date: July, 1988

This is a copy of UCI ML iris datasets.
<http://archive.ics.uci.edu/ml/datasets/Iris>

The famous Iris database, first used by Sir R.A Fisher

This is perhaps the best known database to be found in the pattern recognition literature. Fisher's paper is a classic in the field and is referenced frequently to this day. (See Duda & Hart, for example.) The data set contains 3 classes of 50 instances each, where each class refers to a type of iris plant. One class is linearly separable from the other 2; the latter are NOT linearly separable from each other.

References

- Fisher, R.A. "The use of multiple measurements in taxonomic problems"
Annual Eugenics, 7, Part II, 179-188 (1936); also in "Contributions to Mathematical Statistics" (John Wiley, NY, 1950).
- Duda, R.O., & Hart, P.E. (1973) Pattern Classification

and Scene Analysis.

(Q327.D83) John Wiley & Sons. ISBN 0-471-22361-1.
See page 218.

- Dasarathy, B.V. (1980) "Nosing Around the Neighborhood: A New System

Structure and Classification Rule for Recognition in Partially Exposed

Environments". IEEE Transactions on Pattern Analysis and Machine

Intelligence, Vol. PAMI-2, No. 1, 67-71.

- Gates, G.W. (1972) "The Reduced Nearest Neighbor Rule". IEEE Transactions

on Information Theory, May 1972, 431-433.

- See also: 1988 MLC Proceedings, 54-64. Cheeseman et al's AUTOCLASS II

conceptual clustering system finds 3 classes in the data.

- Many, many more ...

In [154]: **print** iris.data

```
[[ 5.1  3.5  1.4  0.2]
 [ 4.9  3.   1.4  0.2]
 [ 4.7  3.2  1.3  0.2]
 [ 4.6  3.1  1.5  0.2]
 [ 5.   3.6  1.4  0.2]
 [ 5.4  3.9  1.7  0.4]
 [ 4.6  3.4  1.4  0.3]
 [ 5.   3.4  1.5  0.2]
 [ 4.4  2.9  1.4  0.2]
 [ 4.9  3.1  1.5  0.1]
 [ 5.4  3.7  1.5  0.2]
 [ 4.8  3.4  1.6  0.2]
 [ 4.8  3.   1.4  0.1]
 [ 4.3  3.   1.1  0.1]
 [ 5.8  4.   1.2  0.2]
 [ 5.7  4.4  1.5  0.4]
 [ 5.4  3.9  1.3  0.4]
 [ 5.1  3.5  1.4  0.3]
 [ 5.7  3.8  1.7  0.3]
 [ 5.1  3.8  1.5  0.3]
 [ 5.4  3.4  1.7  0.2]
 [ 5.1  3.7  1.5  0.4]
```


[4.6	3.6	1.	0.2]
[5.1	3.3	1.7	0.5]
[4.8	3.4	1.9	0.2]
[5.	3.	1.6	0.2]
[5.	3.4	1.6	0.4]
[5.2	3.5	1.5	0.2]
[5.2	3.4	1.4	0.2]
[4.7	3.2	1.6	0.2]
[4.8	3.1	1.6	0.2]
[5.4	3.4	1.5	0.4]
[5.2	4.1	1.5	0.1]
[5.5	4.2	1.4	0.2]
[4.9	3.1	1.5	0.1]
[5.	3.2	1.2	0.2]
[5.5	3.5	1.3	0.2]
[4.9	3.1	1.5	0.1]
[4.4	3.	1.3	0.2]
[5.1	3.4	1.5	0.2]
[5.	3.5	1.3	0.3]
[4.5	2.3	1.3	0.3]
[4.4	3.2	1.3	0.2]
[5.	3.5	1.6	0.6]
[5.1	3.8	1.9	0.4]
[4.8	3.	1.4	0.3]
[5.1	3.8	1.6	0.2]
[4.6	3.2	1.4	0.2]
[5.3	3.7	1.5	0.2]
[5.	3.3	1.4	0.2]
[7.	3.2	4.7	1.4]
[6.4	3.2	4.5	1.5]
[6.9	3.1	4.9	1.5]
[5.5	2.3	4.	1.3]
[6.5	2.8	4.6	1.5]
[5.7	2.8	4.5	1.3]
[6.3	3.3	4.7	1.6]
[4.9	2.4	3.3	1.]
[6.6	2.9	4.6	1.3]
[5.2	2.7	3.9	1.4]
[5.	2.	3.5	1.]
[5.9	3.	4.2	1.5]
[6.	2.2	4.	1.]
[6.1	2.9	4.7	1.4]
[5.6	2.9	3.6	1.3]
[6.7	3.1	4.4	1.4]
[5.6	3.	4.5	1.5]

[5.8	2.7	4.1	1.]
[6.2	2.2	4.5	1.5]
[5.6	2.5	3.9	1.1]
[5.9	3.2	4.8	1.8]
[6.1	2.8	4.	1.3]
[6.3	2.5	4.9	1.5]
[6.1	2.8	4.7	1.2]
[6.4	2.9	4.3	1.3]
[6.6	3.	4.4	1.4]
[6.8	2.8	4.8	1.4]
[6.7	3.	5.	1.7]
[6.	2.9	4.5	1.5]
[5.7	2.6	3.5	1.]
[5.5	2.4	3.8	1.1]
[5.5	2.4	3.7	1.]
[5.8	2.7	3.9	1.2]
[6.	2.7	5.1	1.6]
[5.4	3.	4.5	1.5]
[6.	3.4	4.5	1.6]
[6.7	3.1	4.7	1.5]
[6.3	2.3	4.4	1.3]
[5.6	3.	4.1	1.3]
[5.5	2.5	4.	1.3]
[5.5	2.6	4.4	1.2]
[6.1	3.	4.6	1.4]
[5.8	2.6	4.	1.2]
[5.	2.3	3.3	1.]
[5.6	2.7	4.2	1.3]
[5.7	3.	4.2	1.2]
[5.7	2.9	4.2	1.3]
[6.2	2.9	4.3	1.3]
[5.1	2.5	3.	1.1]
[5.7	2.8	4.1	1.3]
[6.3	3.3	6.	2.5]
[5.8	2.7	5.1	1.9]
[7.1	3.	5.9	2.1]
[6.3	2.9	5.6	1.8]
[6.5	3.	5.8	2.2]
[7.6	3.	6.6	2.1]
[4.9	2.5	4.5	1.7]
[7.3	2.9	6.3	1.8]
[6.7	2.5	5.8	1.8]
[7.2	3.6	6.1	2.5]
[6.5	3.2	5.1	2.]
[6.4	2.7	5.3	1.9]

```
[ 6.8  3.   5.5  2.1]
[ 5.7  2.5  5.   2. ]
[ 5.8  2.8  5.1  2.4]
[ 6.4  3.2  5.3  2.3]
[ 6.5  3.   5.5  1.8]
[ 7.7  3.8  6.7  2.2]
[ 7.7  2.6  6.9  2.3]
[ 6.   2.2  5.   1.5]
[ 6.9  3.2  5.7  2.3]
[ 5.6  2.8  4.9  2. ]
[ 7.7  2.8  6.7  2. ]
[ 6.3  2.7  4.9  1.8]
[ 6.7  3.3  5.7  2.1]
[ 7.2  3.2  6.   1.8]
[ 6.2  2.8  4.8  1.8]
[ 6.1  3.   4.9  1.8]
[ 6.4  2.8  5.6  2.1]
[ 7.2  3.   5.8  1.6]
[ 7.4  2.8  6.1  1.9]
[ 7.9  3.8  6.4  2. ]
[ 6.4  2.8  5.6  2.2]
[ 6.3  2.8  5.1  1.5]
[ 6.1  2.6  5.6  1.4]
[ 7.7  3.   6.1  2.3]
[ 6.3  3.4  5.6  2.4]
[ 6.4  3.1  5.5  1.8]
[ 6.   3.   4.8  1.8]
[ 6.9  3.1  5.4  2.1]
[ 6.7  3.1  5.6  2.4]
[ 6.9  3.1  5.1  2.3]
[ 5.8  2.7  5.1  1.9]
[ 6.8  3.2  5.9  2.3]
[ 6.7  3.3  5.7  2.5]
[ 6.7  3.   5.2  2.3]
[ 6.3  2.5  5.   1.9]
[ 6.5  3.   5.2  2. ]
[ 6.2  3.4  5.4  2.3]
[ 5.9  3.   5.1  1.8]]
```

```
In [162]: print plt.scatter(iris.data[:,1], iris.data[:,2], c=iris.target)

<matplotlib.collections.PathCollection object at 0x10b8a43d0>
```



Next Steps

Practice doing much of the same functionality we did today with three different data sets from [here \(http://vincentarelbundock.github.io/Rdatasets/datasets.html\)](http://vincentarelbundock.github.io/Rdatasets/datasets.html).

Primarily, your goals are to: * Show how to load a csv file into python * Use pandas to understand the numerical portions of the data * Use matplotlib to visualize the data * Use scikit-learn to fit the data to some dependent feature (y)